

Alex Doe

San Francisco, CA, USA | +1 (555) 123-4567 | alex.doe@example.com | linkedin.com/in/alex-doe-demo | github.com/alex-doe-demo

EDUCATION

State University

M.S. in Computer Science

Anytown, USA

Sep. 2023 – Dec. 2024

Example College

B.S. in Computer Science

Anytown, USA

Sep. 2019 – Jul. 2023

PROFESSIONAL EXPERIENCE

Software Engineer

ExampleCorp

Jan. 2025 – Present

San Francisco, CA

- Built a high-throughput event ingestion service in Go handling ~50K req/s; reduced p99 latency from 220ms to 45ms by introducing batched writes, connection pooling, and a sharded in-memory buffer with backpressure
- Designed and shipped an internal feature-flag SDK adopted by 12 services; implemented gradual rollout, percentage targeting, and a Redis-backed audit log enabling sub-second flag propagation across regions
- Led migration of a legacy monolith auth flow to an OAuth2 + JWT microservice; coordinated with 4 teams, wrote backwards-compatible shims, and rolled out to 100% of users with zero downtime

Software Engineer Intern

DemoLabs

Jun. 2024 – Aug. 2024

Remote

- Prototyped a retrieval-augmented chatbot over internal documentation using LangChain, FAISS, and OpenAI embeddings; achieved 85% answer accuracy on a 200-question eval set, beating keyword-search baseline by 32%
- Optimized vector index ingestion by parallelizing embedding calls and adding incremental updates; reduced full reindex time from 4 hours to 35 minutes on a 200K-document corpus
- Wrote end-to-end tests covering 90% of critical paths; integrated with GitHub Actions to gate deploys, catching 6 regressions before they reached staging

Software Engineer Intern

SampleSoft

Jun. 2023 – Aug. 2023

Anytown, USA

- Implemented a distributed task scheduler in Python with Celery and RabbitMQ; supported priority queues, retries with exponential backoff, and dead-letter routing for ~1.2M daily jobs
- Refactored a 5K-line legacy data-export module into a typed, testable streaming pipeline; raised test coverage from 12% to 78% and cut wall-clock export time by 40%
- Built an internal admin dashboard with React, TypeScript, and Tailwind; reduced manual ops time for the support team by ~6 hours per week

PROJECTS

DistCache – In-Memory Distributed Cache

Oct. 2023 – Feb. 2024

- Built a Redis-compatible distributed cache in C++20 with consistent-hashing partitioning across nodes; benchmarked at 1.2M ops/s on a 3-node cluster with sub-millisecond p99 read latency
- Implemented Raft-based replication for leader election and log-replicated writes across replicas; survived chaos-tested rolling restarts and network partitions with no data loss
- Optimized the hot read path using lock-free hash maps, custom slab allocators, and zero-copy network buffers; reduced p50 GET latency from 120us to 38us

DataFlow – Configurable ETL Pipeline CLI

Apr. 2024 – Jul. 2024

- Designed a YAML-configured ETL CLI on top of DuckDB; supported CSV, Parquet, and Postgres sources with declarative joins, filters, and aggregations executed against an in-memory columnar engine
- Built a typed schema-validation layer with Pydantic that catches type mismatches and missing columns before pipeline execution; reduced data-quality incidents by 6x on a real production workload
- Added structured run logs with per-step row counts, runtime, and memory usage written as JSON to S3; integrated with Datadog dashboards for trend monitoring across ~80 daily pipeline runs

SKILLS

Languages: C++, Go, Python, SQL, Bash

Frameworks & Tools: Raft, gRPC, Protobuf, Docker, Linux, Git, Datadog, AWS